

Lecture 2

Goals

- what is a neural net?
- intuition on structure
- weights, biases + activations
- neural nets as universal approximator of continuous functions

Image Recognition

Trivial for humans \Rightarrow challenging for machines

- can identify different "forms" of the # 3

Computer sees value of every single pixel

- takes in giant matrix \Rightarrow tries to recognize what # it is

Neural net can help us do these types of recognition

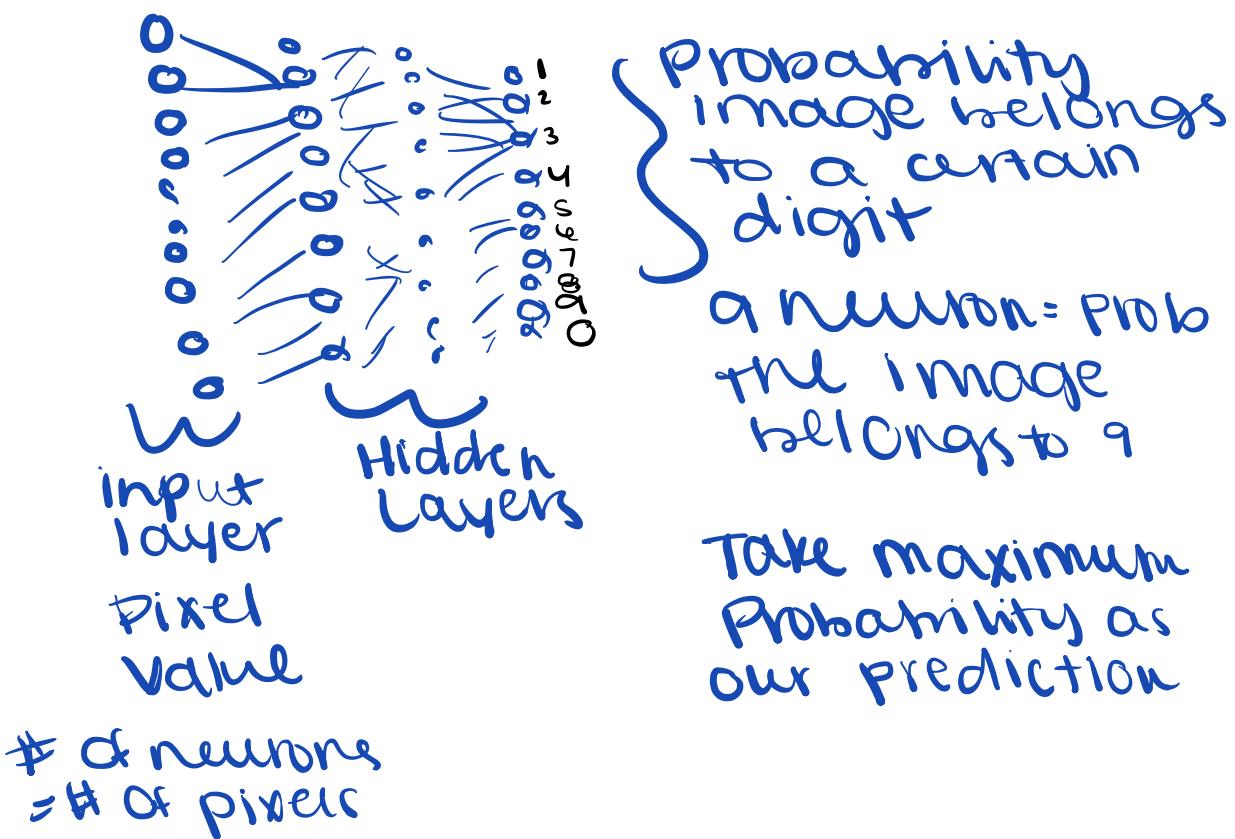
why layers?

- * just talking about plain neural network
 - ↳ multi-layer perceptron

Main ingredients

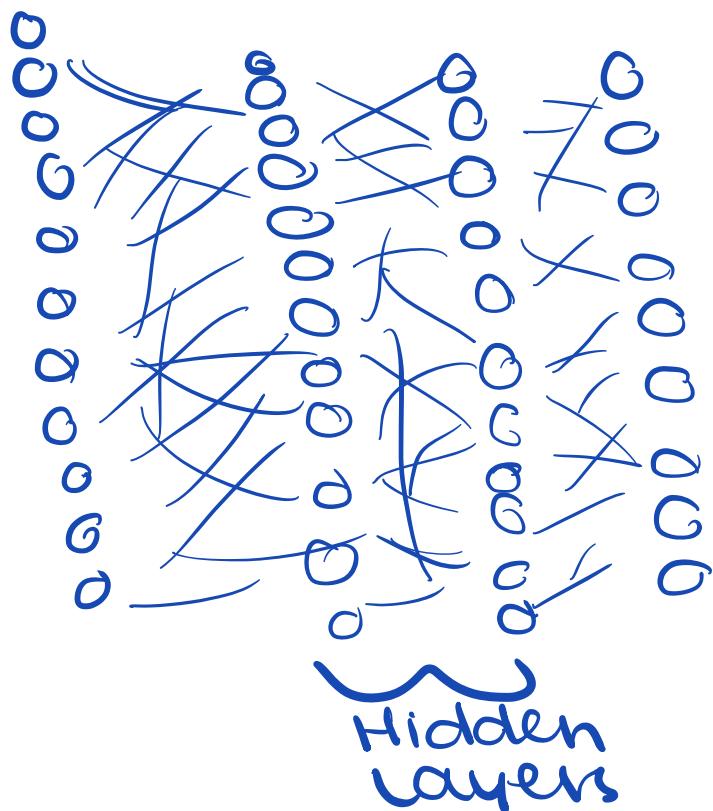
neurons: things that hold a number

↳ takes in each computer pixel as value the neuron holds



Layers used to hold info we care about

NO reverse neuron connection,
all directed



can have any # of neurons +
hidden layers

Dissecting Image Recognition

$$9 \Rightarrow 0 + 1$$

{ looks like a 9
because we see components

$$8 \Rightarrow 0 + 0$$

$4 \Rightarrow | + - . |$

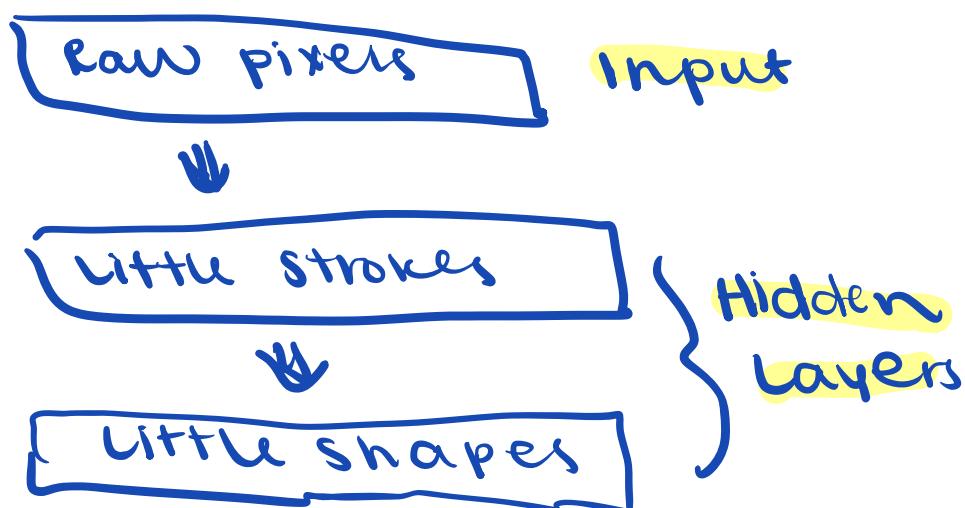
our brain recognizes the digit by looking for essential features in their relative location

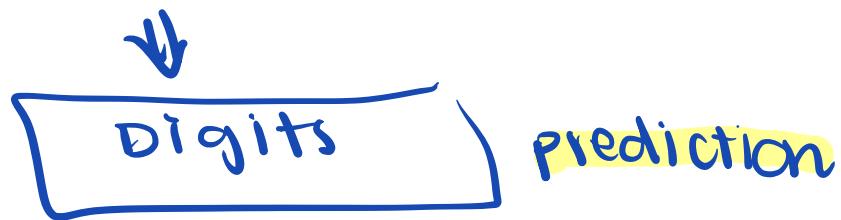
Hidden layers may pick up small patterns in digits to make predictions

we can even further dissect the digit components

$0 \Rightarrow ' + \cup + - +) +)$

this drawing sucks but you get the idea





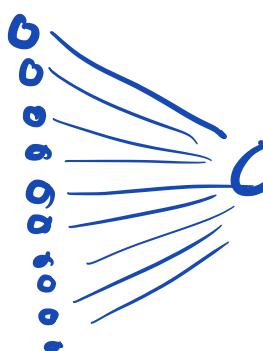
Layers \Rightarrow complex tasks can be decomposed to different levels of granularity

Interpretability is still an open question in AI

- we don't know how to interpret neural nets or if it's using smaller pieces at even single layer

Connections in Neural Nets

passes on info via weights

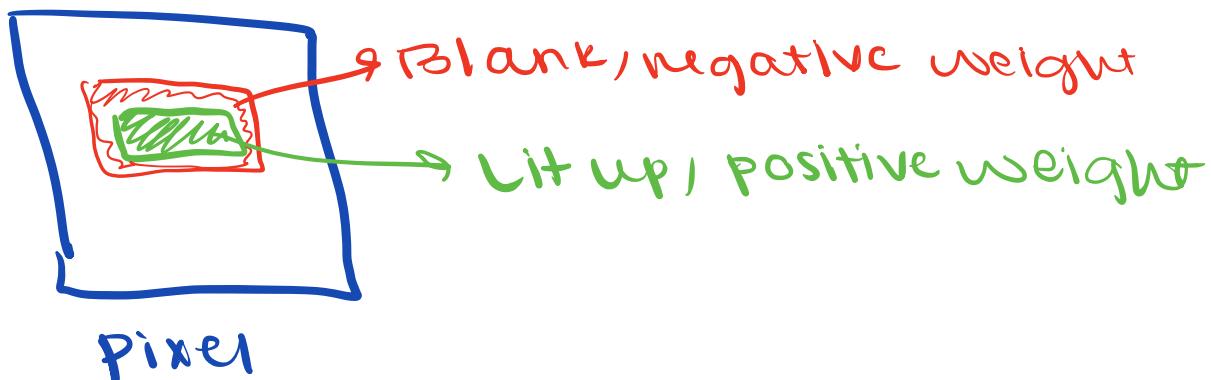


$$\begin{aligned} w_1 &= \# \\ w_2 &= \# \\ w_3 &= \# \\ w_4 &= \# \\ &\vdots \end{aligned}$$

want neuron to light up when the pattern exists

Also need to think about region outside path

Want it to be dark = use negative weights



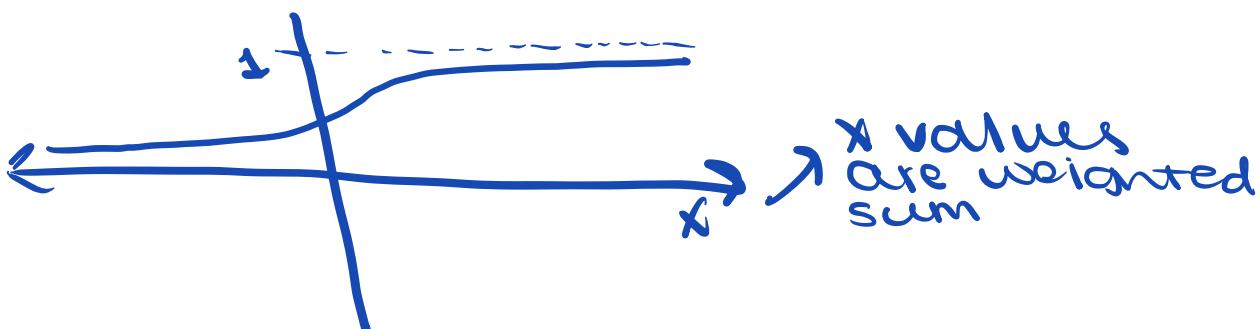
Activation Function

Activation should be in range
0 to 1

Convert weighted combination
to {0,1} w/ activation function

• squash #'s into range

$$O(x) = \frac{1}{1+e^{-x}} \quad] \text{Sigmoid function}$$



Sigmoid lets us capture nonlinear structures

Sigmoid is just one activation function

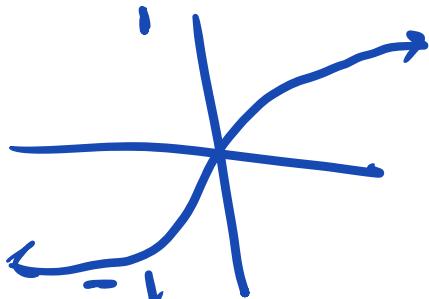
↳ tells us how positive or negative

Tanh (Hyperbolic Tangents)

Maps $-1 \rightarrow 1$

very small $\Rightarrow -1$

very large $\Rightarrow 1$

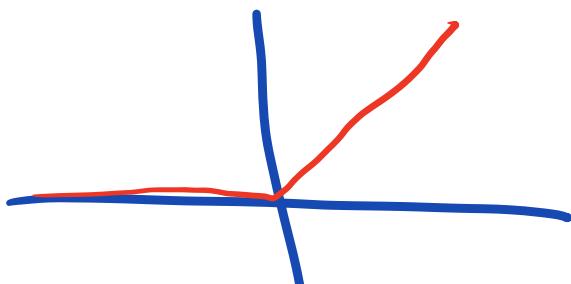


Rectified Linear Unit

$$X \mapsto \max(0, x)$$

- O cuts negative component

- preserves magnitude



|
can activation functions
be linear?

↳ no, makes neural net only
one layer

↳ doesn't add to complexity
or flexibility

Bias in Neural Nets

Bias for inactivity?

$$O(w_0a_0 + w_1a_1 + \dots + w_{n-1}a_{n-1})$$

use bias so
it doesn't
activate when
sum > 0

↓
Bias, only
activates when
weighted sum
is > 10

every neuron has a bias

Learning in Neural Nets

Learning: finding right weights
+ biases for every neuron + layer

Compact Notation

$$a_0' = \sigma(w_{0,0}a_0^{(0)} + w_{0,1}a_1^{(0)} + w_{0,2}a_2^{(0)} + \dots + w_{0,n}a_n^{(0)} + b_0)$$

sigmoid

bias

↳ looking at 1 neuron

a_0 = 0th layer

weighted input + bias

↳ into activation function

we have one equation for
each node



 weights

 one input vector

can rewrite equation using matrix

Product of row of weights
+ column of inputs

$$a_i^{(l)} = w_{i,0}a_0^{(l)} + w_{i,1}a_1^{(l)} + \dots + w_{i,n}a_n^{(l)}$$

different weights, same set of inputs

$$\Theta \left(\begin{bmatrix} w_{00} & w_{01} & \dots & w_{0n} \\ w_{10} & w_{11} & \dots & w_{1n} \\ \vdots & \vdots & \ddots & \vdots \\ w_{K0} & w_{K1} & \dots & w_{Kn} \end{bmatrix} \begin{bmatrix} a_0^{(l)} \\ a_1^{(l)} \\ \vdots \\ a_K^{(l)} \end{bmatrix} \right) \begin{bmatrix} b_0 \\ b_1 \\ \vdots \\ b_K \end{bmatrix}$$

$K \times n$ $n \times 1$ $K \times 1$

$$a' = \Theta(w_a^{(l)} + b)$$

This notation also makes coding easier!

- Languages optimized for vectorized computation

Neurons As Functions

neuron doesn't hold * one number*
↳ different inputs will produce
different outputs

Neural network \approx big complicated
function

Training Neural Nets

start w/ initial net w/ random
values for weights

↳ typically get incorrect
predictions

↳ calculate how far prediction
is from true value

↳ send signal back to update weights

Neural nets as universal Approximator

Can approx. any continuous function w/ a linear comb of translated/scaled ReLU functions

- Single hidden layer MLP
- But may require infinite neurons in the layer

$$f(x) = \sigma(w_1x_1 + b_1) \cdot \sigma(w_2x_2 + b_2) + \dots$$

- Apply weights to even input
+ biases
- Apply sigmoid function to each output

$$\sigma(x) = \max(0, x)$$

- use small segments to approx continuous function

Approx any Boolean, classification or regression function to arbitrary precision

We need to specify how wide

- big we want hidden layers to be
 - Tune # of layers + neurons as iterative process

neurons can ↑ from layer to layer

* Neural net can represent any function if it has sufficient capacity

↳ Not all architecture can represent any function